

LIVE NOW

Stage 2 — Productized B2C

A scalable, self-serve product built on the same assessment IP — no founder hours per user. New territory: nothing in the existing docs prices or scopes this yet.

ROADMAP

KEY PARTNERS - Developer / dev team (see DEVELOPER_HANDOFF selection criteria) - Legal / regulatory counsel - Payment & subscription infrastructure - Later: EMA / passive-sensing vendor (already scouted in docs/research)	KEY ACTIVITIES - Engineering build: typed framework, auth, encrypted persistence, voice pipeline Regulatory review of the no-human-review boundary — gating - Automated report QA at scale (replaces 1:1 founder review — hardest open design problem) - SEO / content engine KEY RESOURCES - The app/platform itself — not yet built - Reusable: assessment IP, scoring engine, Stage 1 brand & trust, de-identified case proof	VALUE PROPOSITIONS Self-serve capacity map: complete the assessment on your own time, get a structured report without booking a \$500 human engagement first. Lower price, lower friction, private. Optional tracking subscription: longitudinal check-ins — the assessment already asks what changes when sleep/timing improve.	CUSTOMER RELATIONSHIPS - Self-serve, automated onboarding - Content / education instead of 1:1 calls - Upsell path back into Stage 1 for complex/high-need cases CHANNELS - SEO/content (keywords already researched in market research doc) - App store / web app direct - Paid search/social at scale, once claims are locked down - Stage 1 email list as seed / cross-sell audience	CUSTOMER SEGMENTS - Same core persona, broader funnel: price-sensitive, earlier-stage, not ready for a human engagement - Wants privacy / self-serve over a sales call
COST STRUCTURE Upfront dev build: placeholder \$30K–\$100K+ (get a real quote). Ongoing: hosting, AI/voice inference per user, payment processing, support, compliance retainer. Marginal cost per user is low — that's the entire point of Stage 2.	REVENUE STREAMS (HYPOTHESES) One-time report ~\$79–\$149 Tracking subscription ~\$15–\$25/mo — bounded by Inflow (\$47.99/mo self-serve app) and Shimmer (~\$260/mo human coaching). Untested — validate with price testing.			

 **Compliance posture changes here.** Stage 1's safety case rests on the founder reviewing every report before a client sees it. Self-serve removes that gate. Get a regulatory read on the automated-report boundary — FDA SaMD exposure, FTC health-claim risk at scale, PMHNP scope-of-practice for an unsupervised consumer product — before this becomes an engineering commitment.

The Bridge: Stage 1 → Stage 2

Stage 1 funds itself and proves the offer today. It also funds and de-risks Stage 2 — if these triggers say it's time.

What carries over

- Brand & trust built during Stage 1 delivery
- The assessment IP and deterministic scoring engine — no rebuild needed
- De-identified case proof and sample outputs
- Email list / warm audience from Stage 1 clients and leads
- Market research already completed (TAM/SAM/SOM, competitor landscape, keyword demand)
- Referral-partner relationships — a Stage 2 channel too

What's net-new

- The engineering build itself (typed framework, auth, encrypted persistence, voice pipeline)
- Self-serve pricing and positioning — untested
- Regulatory posture — no human-review gate by default
- Marketing motion — SEO/paid/app-store vs. referral-led organic
- Automated report QA to replace 1:1 founder review

TRIGGER	METRIC	TARGET
Proof of demand	Cumulative Beta + Sprint + Pilot clients completed	≥ 20
Delivery quality	"Next conversation clarity" client rating	≥ 8/10
Delivery quality	Refund rate	≤ 5%
Capacity ceiling reached	Consecutive months turning away qualified leads at full capacity	≥ 2
Organic pull	Share of new leads arriving without founder outreach	≥ 25%
Price/format objection signal	Share of fit-call leads citing self-serve or price friction	≥ 15%
Compliance groundwork	Security threat model + regulatory review of the no-human-review boundary complete	Yes
Funding secured	Stage 1 profit reserve or capital covers the dev build quote	Yes

 Live tracker with formulas: **Voice_EC_Map_Unit_Economics.xlsx**, tab **Transition_Triggers**. Fill in "Current Status" as real numbers come in — the sheet flags a recommendation automatically once enough triggers are met.